Visibility Is Not Detection: Auditing Measurement Boundaries in Staged and Evidence-Ranking Evaluation

Author TBD_OWNER | Affiliation TBD_OWNER | Corresponding author TBD_OWNER

Structured abstract

Purpose. We audit whether reported evaluation scores distinguish observed defects from silent absorption and whether evidence-ranking conclusions survive implementation choices. Methods. We retain the original two-case scope: a frozen staged document evaluator with controlled header perturbations, instrumentation and a nested FinQA measurement-object probe; and FinRank, a public provenance-sensitive evidence-ranking benchmark. Results. In the staged case, 96/118 header injections are visible, the OCR-like family is unobserved on the frozen surface (0/13), and flat versus layered scores are affine/rank-equivalent. In FinRank, tie policy changes absolute metrics but preserves the hard-versus-random direction; a post-hoc text-sha1 correction finds 442 overlapping hard-negative occurrences across 305 records and gold loss in 33 records. Conclusion. The two cases support a bounded review discipline: bind observations to measurement objects, separate score sensitivity from decision stability, preserve identity contracts, and report unresolved reproduction gaps.

Reviewer-pass scope

The contribution is a six-object measurement discipline instantiated in two existing audits. The FinRank identity correction is post-hoc v2 clustered by canonical text_sha1 and is not a second frozen application. OCR-like invisibility, deployment effects of alternative tie policies, and global/metadata metric drift remain unresolved. This package does not claim acceptance, publication, school recognition, or C2/C3 experiments.

Owner fields

All author, affiliation, corresponding-author, funding, contribution, signature, payment, and public-release fields remain TBD_OWNER actions.